

# MCMC Sampling

Lennox Garay

2023-10-16

```
library(boot)
library(tidyverse)

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.2      v readr      2.1.4
## v forcats    1.0.0      v stringr   1.5.0
## v ggplot2    3.4.3      v tibble    3.2.1
## v lubridate  1.9.2      v tidyr     1.3.0
## v purrr      1.0.1
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors

library(MASS)

##
## Attaching package: 'MASS'
##
## The following object is masked from 'package:dplyr':
##
##      select

library(mvtnorm)

set.seed(538)
mu <- matrix(c(0,10), ncol=1)
sig <- matrix(c(4,10,10,100), nrow=2, byrow = TRUE)

post <- function(alpha,beta){
  x = c(-0.86, -0.3, -0.05, 0.73)
  n <- c(5, 5, 5, 5)
  y <- c(0, 1, 3, 5)
  likelihood = invlogit(alpha + beta * x)
  p = prod(likelihood^y * (1 - likelihood)^(n - y))*dmvnorm(c(alpha, beta), as.vector(mu), sig)
  return(p)
}

post = function(alpha, beta) {
  x = c(-0.86, -0.3, -0.05, 0.73) #the data
  n = c(5, 5, 5, 5)
  y = c(0, 1, 3, 5)
  alpha.prior = dnorm(alpha, mean = 0, sd = 2)
```

```

beta.prior = dnorm(beta, mean = 10, sd = 10)
temp = inv.logit(alpha + beta * x)
p = prod(temp^y * (1 - temp)^(n - y))*alpha.prior*beta.prior
return(p)
}

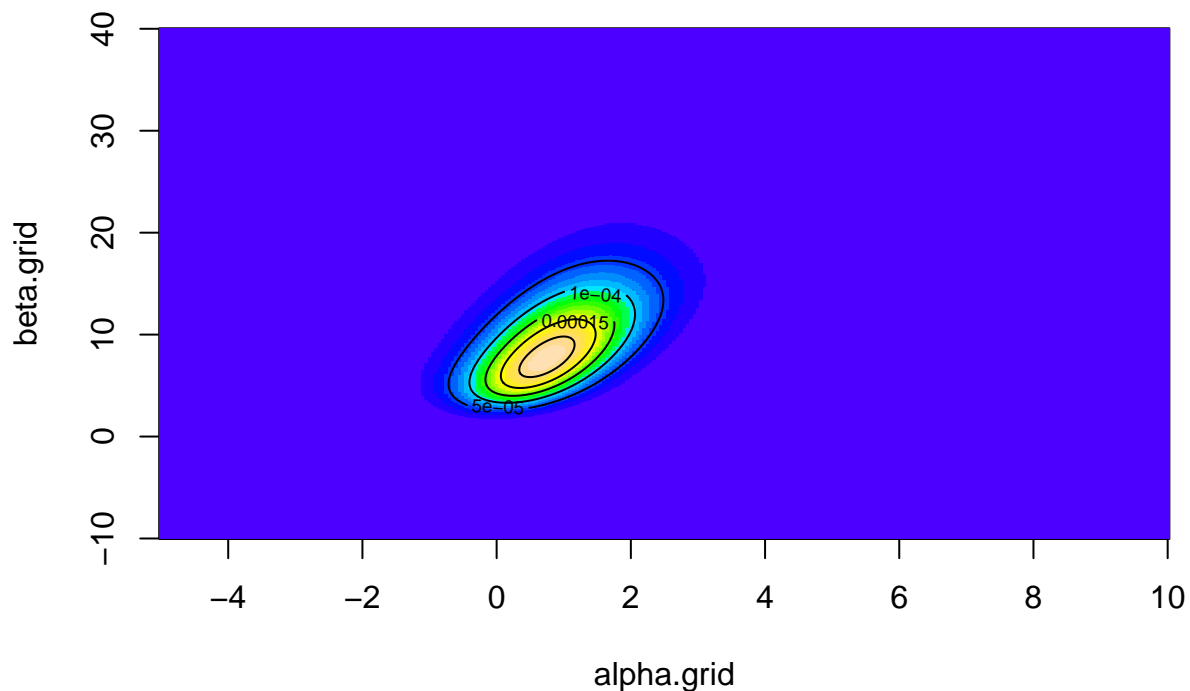
alpha.grid <- seq(-5, 10, 0.05)
a.n <- length(alpha.grid)
beta.grid <- seq(-10, 40, 0.1)
b.n <- length(beta.grid)
post.grid <- matrix(0, nrow = a.n, ncol = b.n)

alpha.post = seq(0, length.out=a.n)
for (i in 1:a.n) {
  for (j in 1:b.n) {
    post.grid[i, j] <- post(alpha.grid[i], beta.grid[j])
  }

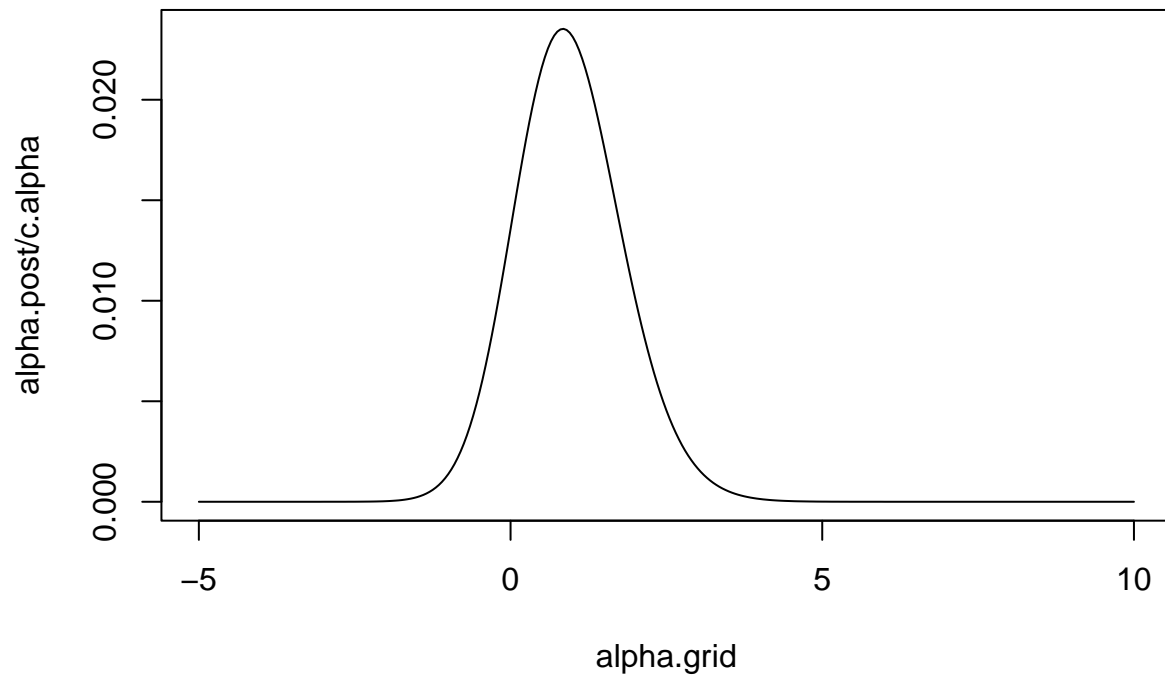
  alpha.post[i] <- sum(post.grid[i, ])
}
c.joint <- sum(post.grid) #Finding the normalizing constants
c.alpha <- sum(alpha.post)

image(alpha.grid, beta.grid, post.grid/c.joint, col = topo.colors(20))
contour(alpha.grid, beta.grid, post.grid/c.joint,
add = TRUE)

```



```
plot(alpha.grid, alpha.post/c.alpha, type = "l")
```




1 (a): Inverse CDF Sampling

```
n = 1000
U = runif(n, 0,1)
```

```
iweib = function(U,alpha,beta){
  ex = (1/alpha)
  temp = log(1-U)
  q = ((-1*(temp))ex)* beta
  return(q)
}
```

```
Fx = iweib(U,alpha=2, beta=0.5)
truth = rweibull(n,2,0.5)
grid = seq(-50,50, length.out=1000)
df = data.frame(U,Fx, truth,grid)
```

```
ggplot(df, aes(x=Fx)) + geom_density(aes(x=Fx), colour='gold', lwd=1.5) + geom_density(aes(x=truth), co
```



1(b):

```
alpha = 2
beta = 0.5
k = alpha
N = 10000
U = runif(N,0,1)
Y = rexp(N, rate=1/beta)
fy = dweibull(Y,alpha,scale=beta) ## "f" in f/g
g = k*dexp(Y,rate=1/beta) #this is our envelope
accepted = Y[U <= (fy/g)]
```


```
length(accepted)/N ## %accepted
```

```
## [1] 0.5037
```

```
theta.grid = seq(0,2, length.out=1000)
```

```
weibb = dweibull(theta.grid,alpha,scale=beta)
```

```
plot(theta.grid,weibb, type='l', col='blue')
lines(density(accepted), type='l', col='red')
```



## 2 (a):

Algorithm:

- 1) Simulate some data from normal (this generates  $y_i$ ).
- 2) Calculate sample variance and sample mean from  $n$  data.
- 3) We draw  $N$  random samples from  $\sigma^2 | y \sim Sc.Inv.\chi^2(n-1, s^2)$
- 4) We know that  $\mu | \sigma^2, y \sim N(\bar{y}, \frac{\sigma^2}{n})$ . (of course this relies on the sampled values of  $\sigma^2$  in step 3.)

MC Integration


```
n = 50
y = rnorm(n, 2, 1)

s = sd(y)
s2 = s^2
ybar = mean(y)

N = 2000
sig2 = rinvchisq(N, df=n-1, scale=s2)
mu = rnorm(N, mean=ybar, sd = sqrt(sig2/n))

post = data.frame(mu, sig2)

ggplot(data=post, aes(x=mu, y=sig2)) + geom_point() + geom_density_2d(colour='orange') + ggtitle('Joint')
```




```
df2 = data.frame(sig2)


griddy = seq(0,3, length.out=length(sig2))

hist(sig2, breaks = 30, prob=T, main = 'MC Samples of Sigma^2')
lines(density(sig2), type='l', col = 'blue', lwd = 1.5)
```

## MC Samples of $\text{Sigma}^2$



```
hist(mu, breaks=30, prob=T, main = 'MC Samples of mu')  
lines(density(mu), type='l', col='blue', lwd=1.5)
```



```
Mode <- function(x) { ## Taken off the web; tried an R package but it gave me the wrong answer.
  ux <- unique(x)
  ux[which.max(tabulate(match(x, ux)))]
}
```

```
cbind(mu=Mode(mu),sig=Mode(sig2))
```

```
##           mu      sig
## [1,] 2.405056 1.204039
```

We see that the MAP is pretty close to the true parameter values! How fun.

## 2 (d):

We have that the posterior predictive is:

$$\int \int p(\tilde{y}|\mu, \sigma^2)p(\mu, \sigma^2|y)d\mu d\sigma^2$$

Which we can sample from a Normal distribution using the posterior samples of  $\mu$  and  $\sigma^2$  so,  $y \sim N(\mu_{post}, \sigma_{post}^2)$

*## using these for the posterior predictive*

```
sig.new = sig2
```

```
mu.new = mu
```

```
y.pred = as.numeric()
```

```
for(i in 1:length(sig.new)){
```

```
  y.pred[i] = rnorm(1, mean=mu.new[i], sd=sqrt(sig.new[i]))
```

```
}
```




```
## is it possible to do this without a for loop? eh lets try

samps = length(sig2)

y.pred2 = rnorm(n=samps, mean=mu.new, sd=sqrt(sig.new))

plot(density(y.pred), col = 'blue', main = 'Posterior Predictive for mu and sigma', lty='dashed')
lines(density(y.pred2), col = 'gold', main = 'Posterior Predictive for mu and sigma?')
```

## Posterior Predictive for mu and sigma



N = 2000 Bandwidth = 0.2175


```
## BEAUTIFUL! NO FOR LOOP NEEDED. We love R vectorization.

## Theoretical Post Pred

sig.real = var(y)

y_samps2 = rt.scaled(2000, df = n-1, mean = ybar, sd = (sqrt(1 + 1/n))*sqrt(sig.real))

hist(y.pred2, breaks=30, main='Post. Pred. Truth in red', prob=T)
lines(density(y_samps2), col = 'red', lty='dashed', lwd=1.5)
```



### E: Gibbs(?!)

```

B = 2000

mu.gib = rep(0, length(y))
sig2.gib = rep(0, length(y))


mu.gib[1] = ybar ## initialization

for(b in 2:B){
  scaled = sum((y - mu.gib[b-1])^2)/n
  sig2.gib[b-1] = rinvchisq(1,df=n,scale=scaled)
  mu.gib[b] = rnorm(1,ybar, sd=sqrt(sig2.gib[b-1]/n))
}


par(mfrow=c(1,2))
plot(sig2.gib, type="l", main = "Trace Plot of Sigma2")
plot(mu.gib, type = "l", main = "Trace Plot of Mu")

```

**Trace Plot of Sigma2**




**Trace Plot of Mu**




```
hist(sig2.gib, breaks=30, main='Gibbs sample of sig^2 post', prob=T)
lines(density(sig2.gib), col='red')
hist(mu.gib, breaks=30, main='Gibbs sample of mu post', prob=T)
lines(density(mu.gib), col='red')
```

**Gibbs sample of  $\sigma^2$  post**



**Gibbs sample of  $\mu$  post**



We were in fact able to use Gibbs sampling on this data. Both MC and Gibbs sampling were pretty good at estimating the true posterior density. It was pretty straight forward, but I think that if MC samples works, then it is best to use that. MC sampling is even simpler than Gibbs.